



rahulvg / MLOPS-Assignment-Group-41-



Code

Issues

Pull requests 1

Actions

Projects

Wiki

Security



mlops\_assignment ▾

MLOPS-Assignment-Group-41- / readme.md



rahulvg Update readme.md

cc09eb1 · now



280 lines (176 loc) • 7.4 KB

# Heart Disease Prediction System – MLOps

**Course:** MLOps

**Assignment:** End-to-End MLOps Pipeline

**Dataset:** UCI Heart Disease Dataset

**Group:** Group 41

**Repository:** <https://github.com/rahulvg/MLOPS-Assignment-Group-41->

## Problem Statement

The objective of this project is to design, develop, and deploy a scalable, reproducible, and production-ready machine learning system to predict the presence of heart disease based on patient health attributes.

The solution follows modern MLOps best practices, including experiment tracking, CI/CD automation, containerization, Kubernetes deployment, and monitoring.

## 1. Setup and Installation Instructions

### 1.1 Local Environment

**Python Version:** 3.10

## Install Dependencies

```
pip install -r requirements.txt
```



## Launch MLflow UI (Local SQLite DB)

```
mlflow ui --backend-store-uri sqlite:///mlflow.db
```



Access MLflow at:

<http://localhost:5000>



## 1.2 Verification of Docker Build and Execution via GitHub Actions

Due to organizational restrictions that prevent local installation of Docker Desktop, the Docker image build and container execution were verified using **GitHub Actions**, which provides a Docker-enabled runner environment.

This ensures that containerization and execution are **reproducible, verifiable, and independent of local system constraints**.

1. Navigate to the GitHub repository:

<https://github.com/rahulvg/MLOPS-Assignment-Group-41->

2. Click on the **Actions** tab in the repository.

3. Select the most recent workflow run under the **CI pipeline**.

4. Open the workflow run and inspect the following steps:

- **Build Docker image**

This step executes the Docker build command using the project's `Dockerfile` .

- **Run Docker container and test API**

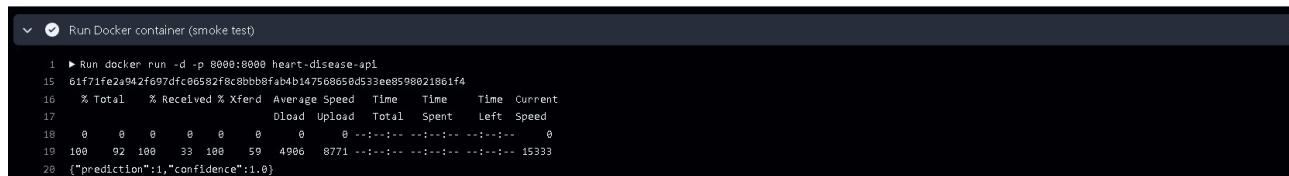
This step starts the container and invokes the `/predict` endpoint using a sample JSON request.

# Evidence of Successful Docker Execution

Within the GitHub Actions workflow logs, the following evidence can be observed:

- Docker build logs confirming successful image creation
- Container startup logs indicating the FastAPI service is running
- Successful HTTP response from the `/predict` endpoint returning a prediction and confidence score

## Screenshot of successfull Docker run event

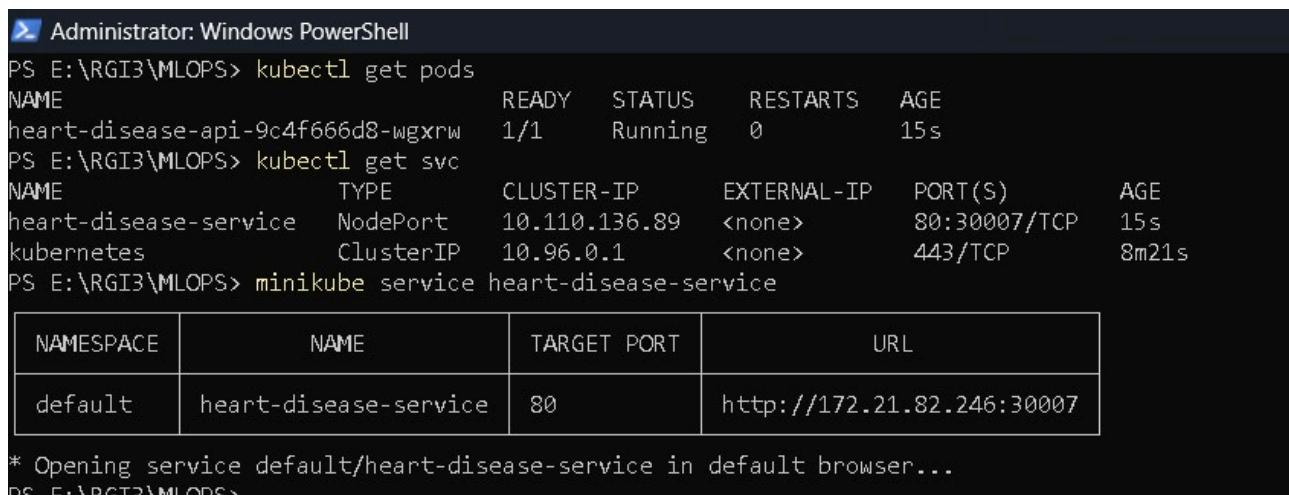


A screenshot of a terminal window showing the output of a Docker command. The command is `Run docker run -d -p 8000:8000 heart-disease-api`. The log output shows the container ID (61f71fe2a942fc097dfc06582fc8bbb0fa4b147568658d533ee8598021961f4), followed by a table of metrics: Total, Received, Xferd, Average Speed, Time, Time, Current, Download, Upload, Total, Spent, Left, Speed. The last few lines show the API's response: {"prediction":1,"confidence":1.0}.

## 1.3 Kubernetes (Local Deployment with Minikube)

### Start Minikube

```
minikube start --container-runtime=containerd
```



A screenshot of a Windows PowerShell window with administrator privileges. It shows the results of several commands: `kubectl get pods` (listing a pod named `heart-disease-api-9c4f666d8-wgxrw`), `kubectl get svc` (listing a service named `heart-disease-service` with NodePort 80:30007/TCP and ClusterIP 10.96.0.1), and `minikube service heart-disease-service` (displaying a table of service details). The table shows the service is in the default namespace, name is `heart-disease-service`, target port is 80, and the URL is `http://172.21.82.246:30007`.

\* Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

Opening service default/heart-disease-service in default browser...

PS E:\RGIT\MLOPS>

\*

## Deploy Application

```
kubectl apply -f k8s/deployment.yaml  
kubectl apply -f k8s/service.yaml
```

# Expose Service

```
minikube service heart-disease-service
```

 mlops\_assignment ▾ MLOPS-Assignment-Group-41- / readme.md  Top

Preview Code Blame

MLOPS-Assignment-Group-41- / README.md

↑ Top

## 2. Data Acquisition and Exploratory Data Analysis

## 2.1 Dataset

- Source: UCI Machine Learning Repository
  - Format: CSV
  - Task: Binary classification (presence or absence of heart disease)

## 2.2 Preprocessing

- Missing values handled
  - Numerical features scaled using StandardScaler
  - Target variable encoded

- Preprocessing implemented using a scikit-learn Pipeline

## 2.3 Exploratory Data Analysis (EDA) & Modelling choice

- Feature distributions analyzed using histograms
- Correlation heatmap used to study feature relationships
- Class balance verified

The modelling approach was guided by dataset characteristics, interpretability needs, and deployment stability.

Two models were evaluated:

- **Logistic Regression** – chosen as a strong, interpretable baseline for structured medical data
- **Random Forest** – included to capture non-linear relationships and feature interactions

All numerical features were standardized using **StandardScaler**, and preprocessing was implemented through a unified **scikit-learn Pipeline** to ensure reproducibility, prevent data leakage, and enable deployment-safe inference.

### Hyperparameter Tuning

- Logistic Regression:  $c \in \{0.1, 1.0, 10.0\}$
- Random Forest:
  - $n\_estimators \in \{100, 200\}$
  - $max\_depth \in \{None, 10\}$

Each configuration was logged as a separate experiment using **MLflow**.

### Evaluation

Models were evaluated using **5-fold cross-validation** with the following metrics:

- Accuracy
- Precision
- Recall
- ROC-AUC

### Final Model

**Logistic Regression with C = 0.1** was selected due to:

- Consistent cross-validation performance
- Lower variance across folds
- Better generalization
- Simpler and more interpretable behavior

Its stability and ease of monitoring make it well-suited for a production-oriented MLOps pipeline.

### 3. Experiment Tracking

MLflow was integrated to track:

- Model parameters
- Cross-validation metrics
- Model artifacts

**All experiments are logged under a dedicated MLflow experiment for easy comparison.**

Run Name	Created	Dataset	Duration	Source	Models
Final_LR_C0.1	1 day ago	-	3.0s	C:\Users...	model
Final_LR_C0.1	1 day ago	-	3.0s	C:\Users...	model
Final_LR_C0.1	1 day ago	-	4.1s	C:\Users...	model
Final_LR_C0.1	1 day ago	-	3.0s	C:\Users...	model
Final_LR_C0.1	1 day ago	-	3.0s	C:\Users...	model
Final_LR_C0.1	1 day ago	-	4.1s	C:\Users...	model
Final_LR_C0.1	1 day ago	-	8.4s	E:\ML...\	model
RF_ne200_depth10	1 day ago	-	2.9s	E:\ML...\	-
RF_ne200_depthNone	1 day ago	-	3.0s	E:\ML...\	-
RF_ne100_depthNone	1 day ago	-	1.8s	E:\ML...\	-
LR_C10.0	1 day ago	-	287ms	E:\ML...\	-
LR_C1.0	1 day ago	-	257ms	E:\ML...\	-
LR_C0.1	1 day ago	-	314ms	E:\ML...\	-
Final_LR_C0.1	1 day ago	-	8.5s	E:\ML...\	model
RF_ne200_depth10	1 day ago	-	3.0s	E:\ML...\	-
RF_ne200_depthNone	1 day ago	-	3.1s	E:\ML...\	-
RF_ne100_depthNone	1 day ago	-	1.9s	E:\ML...\	-
LR_C10.0	1 day ago	-	244ms	E:\ML...\	-
LR_C1.0	1 day ago	-	266ms	E:\ML...\	-

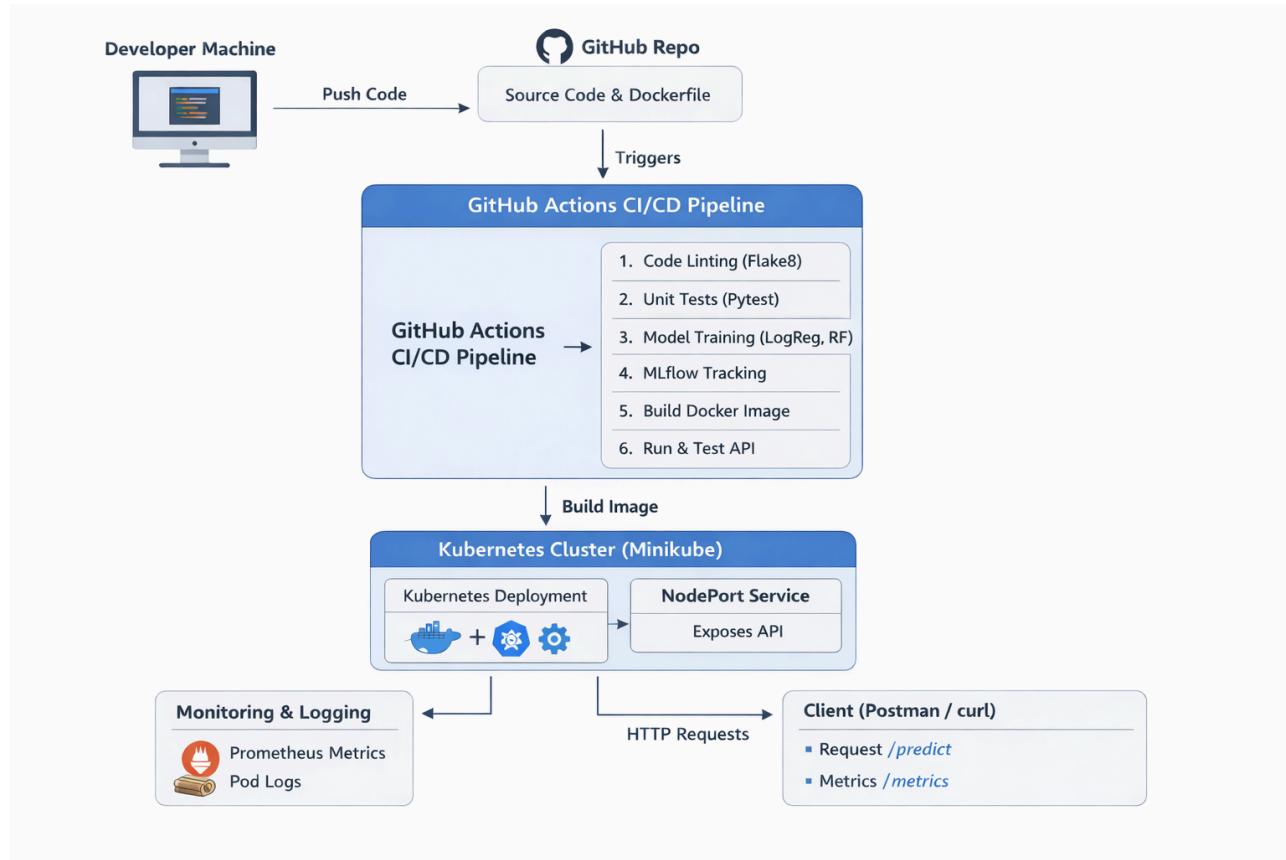
20 matching runs

### Model Packaging and Reproducibility

- Final model saved as a serialized scikit-learn Pipeline
- Model can be found at **final\_model\heart\_disease\_lr\_c01.pkl** in git repo.
- Preprocessing included within the model

- Reproducible inference guaranteed
- Dependencies listed in requirements.txt
- Artifacts stored and versioned using MLflow check `mlflow_experiment.db` in git repo

## 4. Architecture Diagram



## 5. CI/CD Pipeline

### Tools Used

- GitHub Actions
- Pytest
- Flake8
- Docker

### Pipeline Stages

- Code linting

- Unit testing
- Model training
- Docker image build
- API smoke testing

The screenshot shows the GitHub Actions interface for the repository "rahulvg / MLOPS-Assignment-Group-41". The "Actions" tab is selected. On the left, a sidebar lists management options like Caches, Attestations, Runners, Usage metrics, and Performance metrics. The main area displays a table of "23 workflow runs" from various workflows. The columns include Event (e.g., Docker, Merge pull request), Status (green checkmark), Branch (mlops\_assignment, main, init-commit), Actor (rahulvg), Date (Dec 31, 2025, Dec 30, 2025), and Duration (1m 46s to 1m 3s). A search bar at the top right allows filtering by workflow runs.

The screenshot shows the GitHub Actions job logs for the "build-test-train" job. The job summary indicates it succeeded yesterday in 1m 44s. The job details page lists the steps: Checkout code, Set up Python, Install dependencies, Run linting (flake8), Run unit tests, Upload Pytest HTML report, and Train final model. The "Train final model" step is expanded, showing a log of command-line output for a logistic regression experiment. The log includes commands like "Run export PYTHONPATH=\$PWD", "Run export PYTHONPATH=\$PWD", and "Run export PYTHONPATH=\$PWD". It also shows statistics for accuracy, precision, recall, and ROC\_AUC across multiple runs (LR\_C0.1, LR\_C1.0, LR\_C10.0).

## 6. Code Repository

<https://github.com/rahulvg/MLOPS-Assignment-Group-41->

## 7. Containerization and Deployment

---

### 7.1 Dockerized API

- FastAPI-based service
- /predict endpoint
- Accepts JSON input
- Returns prediction and confidence score

### 7.2 Kubernetes Deployment

- Local Kubernetes using Minikube
- Deployment and NodePort Service manifests
- API tested using curl and Postman

## 7. Monitoring and Logging

---

### 7.1 Logging

- Request-level logging implemented via FastAPI middleware
- Logs include endpoint, HTTP status, and latency
- Logs accessible via Kubernetes pod logs

### 7.2 Monitoring

- Prometheus-compatible /metrics endpoint exposed
- Metrics include request count and request latency
- Ready for Prometheus and Grafana integration

```
[Administrator: Windows PowerShell]
PS E:\RGIS\MLDF5> kubectl logs heart-disease-api-9c4f666d8-f998 1/1 Running 0 2m19s
/usr/local/lib/python3.10/site-packages/sklearn/base.py:348: InconsistentVersionWarning: Trying to unpickle estimator StandardScaler from version 1.8.0 when using version 1.3.2. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
  warnings.warn(
/usr/local/lib/python3.10/site-packages/sklearn/base.py:348: InconsistentVersionWarning: Trying to unpickle estimator LogisticRegression from version 1.8.0 when using version 1.3.2. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
  warnings.warn(
/usr/local/lib/python3.10/site-packages/sklearn/base.py:348: InconsistentVersionWarning: Trying to unpickle estimator Pipeline from version 1.8.0 when using version 1.3.2. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
  warnings.warn(
INFO: Started server process [1]
INFO: Waiting for application startup...
INFO: Application startup complete
INFO: Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)
/usr/local/lib/python3.10/site-packages/sklearn/base.py:465: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
/usr/local/lib/python3.10/site-packages/sklearn/base.py:465: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
025-12-31 14:36:55,174 | INFO | POST /predict | status=200 | latency=0.003s
http://127.0.0.1:8000 - POST /predict?client-type=application/json HTTP/1.1" 200 OK
PS E:\RGIS\MLDF5>
```

## 8. Architecture Overview

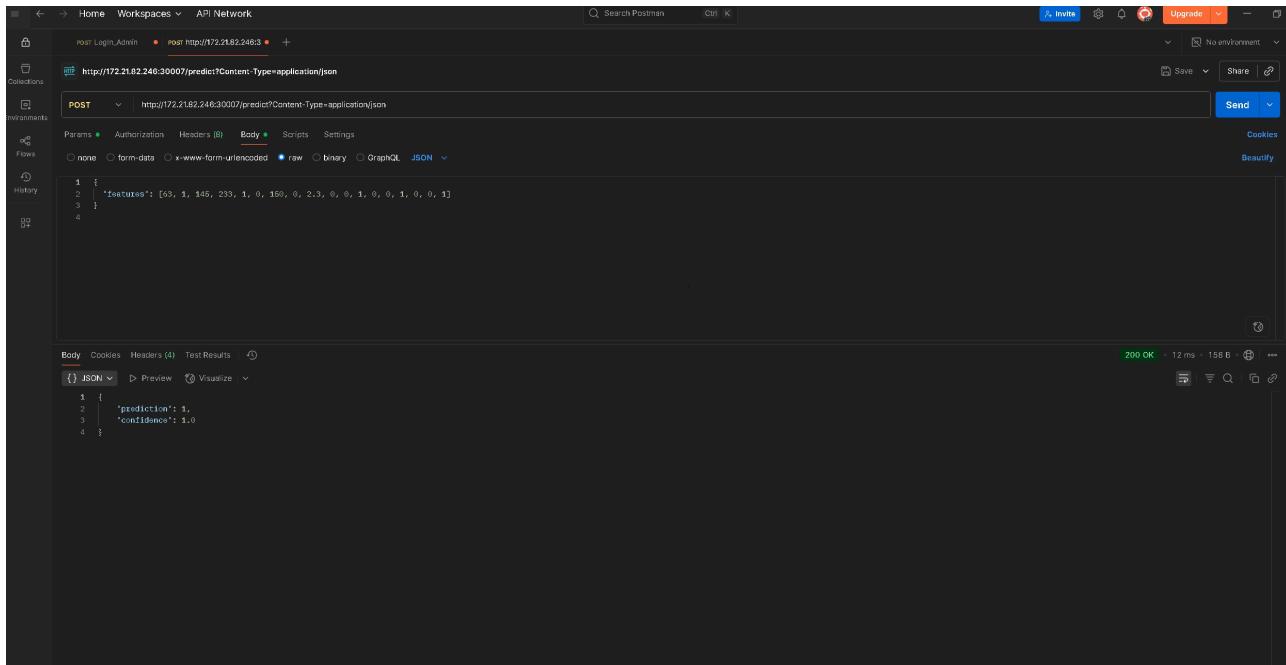
---

Client (Postman / curl)

- FastAPI API (/predict, /metrics)
- Scikit-learn Pipeline
- Kubernetes Pod
- NodePort Service

CI/CD is handled using GitHub Actions, and experiment tracking is handled using MLflow.

# 9 Run using CURL/Postman API



## Conclusion

This project demonstrates a complete, production-grade MLOps workflow covering data analysis, model development, experiment tracking, CI/CD automation, containerization, Kubernetes deployment, and monitoring.

The system is scalable, reproducible, and aligned with real-world MLOps practices.

## Appendix: Useful Commands

### Launch MLflow with Custom Local DB

```
mlflow ui --backend-store-uri sqlite:///E:/RGI3/MLOPS/mlflow.db
```



### Rebuild and Redeploy on Minikube

```
minikube image build -t heart-disease-api .
kubectl delete deployment heart-disease-api
kubectl apply -f k8s/deployment.yaml
```



